

# Speeding up parameter inference with compiled models

*Philipp H Boersch-Supan, Sadie J Ryan, and Leah R Johnson*

*September 12, 2016*

Differential equation (DE) models in R are easy to implement and allow simple interactive development using readable code and access to R's many high-level functions. This was part of our motivation to develop **deBInfer** as an R based inference package. However, numerically solving DE models specified as R functions is also relatively slow. The **deSolve** package also allows the evaluation of DE models that are defined in lower-level languages such as C and FORTRAN. These compiled models have the benefit of increased simulation speed. As the DE model is evaluated many times during the MCMC procedure, even moderate speed-ups from using compiled models can result in large absolute time savings.

We demonstrate the speed-up using an example from the **deSolve** documentation (Soetaert *et al.* 2010). Full details on the model specification can be found in (Soetaert *et al.* 2009) which can be displayed with the command `vignette("compiledCode")`. Further details on the set up of the **deBInfer** inference procedure are described in (Boersch-Supan & Johnson 2016) and annotated examples are available in the vignettes `vignette("logistic_ode_example")` and `vignette("vignette_chytrid_dede_example")`.

## 1 Specifying the ODE model

Following Soetaert et al. (2009), we use the following simple ODE:

$$\begin{aligned}\frac{dy_1}{dt} &= -k_1 \cdot y_1 + k_2 \cdot y_2 \cdot y_3 \\ \frac{dy_2}{dt} &= k_1 \cdot y_1 - k_2 \cdot y_2 \cdot y_3 - k_3 \cdot y_2 \cdot y_2 \\ \frac{dy_3}{dt} &= k_3 \cdot y_2 \cdot y_2\end{aligned}$$

where  $y_1$ ,  $y_2$  and  $y_3$  are state variables, and  $k_1$ ,  $k_2$  and  $k_3$  are parameters.

### 1.1 ODE model implementation in R

We implement this model as an R function:

```
modelR <- function(t, Y, parameters) {  
  with (as.list(parameters), {  
  
    dy1 = -k1*Y[1] + k2*Y[2]*Y[3]  
    dy3 = k3*Y[2]*Y[2]  
    dy2 = -dy1 - dy3  
  
    list(c(dy1, dy2, dy3))  
  })  
}
```

And also specify the Jacobian ( $\frac{\partial y'}{\partial y}$ ) for it:

```

jacR <- function (t, Y, parameters) {
  with (as.list(parameters),{

    PD[1,1] <- -k1
    PD[1,2] <- k2*Y[3]
    PD[1,3] <- k2*Y[2]
    PD[2,1] <- k1
    PD[2,3] <- -PD[1,3]
    PD[3,2] <- k3*Y[2]
    PD[2,2] <- -PD[1,2] - PD[3,2]

    return(PD)
  })
}

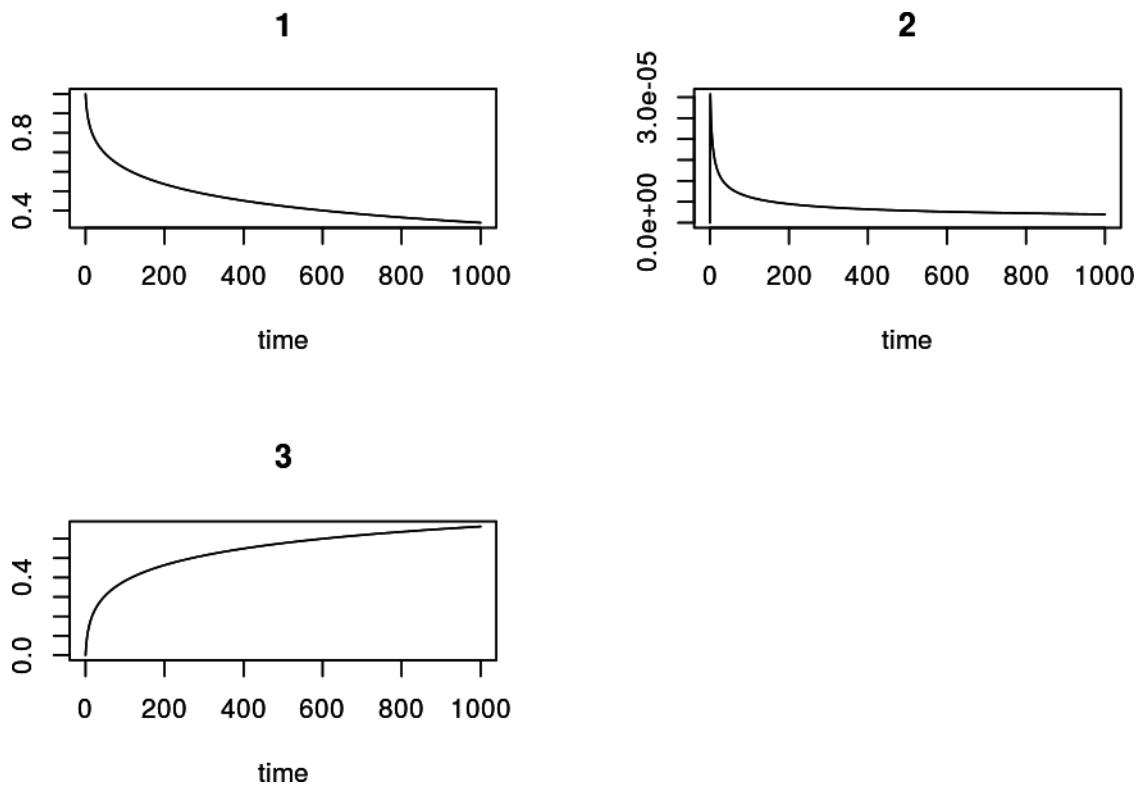
```

This model can then be run as follows:

```

library(deSolve)
parms <- c(k1 = 0.04, k2 = 1e4, k3=3e7)
Y      <- c(1.0, 0.0, 0.0)
times  <- seq(0, 0.1*10^4, length.out = 1000)
PD      <- matrix(nrow = 3, ncol = 3, data = 0)
outR    <- ode(Y, times, modelR, parms = parms, jacfunc = jacR)
plot(outR)

```



## 1.2 ODE model implementation in C

To create compiled models (dynamic-link libraries (.dll) on Windows or shared objects (.so) on other systems) a recent version of the GNU compiler suite is required. This is usually the case for Linux systems. Windows users can install the required toolchain by following the instructions on [\[https://cran.r-project.org/bin/windows/Rtools/\]](https://cran.r-project.org/bin/windows/Rtools/). OSX users need to download and install an appropriate version of Xcode from the Apple developer website [\[https://developer.apple.com/\]](https://developer.apple.com/) or the OSX App Store.

The call to the derivative and Jacobian function is more complex for compiled code compared to R-code, because it has to comply with the interface needed by the integrator. The requirements for this are detailed in Soetaert et al. (2009). A C implementation of the example model is found below.

```
/* file mymod.c */
#include <R.h>
static double parms[3];
#define k1 parms[0]
#define k2 parms[1]
#define k3 parms[2]

/* initializer */
void initmod(void (* odeparms)(int *, double *))
{
    int N=3;
    odeparms(&N, parms);
}

/* Derivatives and 1 output variable */
void derivs (int *neq, double *t, double *y, double *ydot,
             double *yout, int *ip)
{
    if (ip[0] <1) error("nout should be at least 1");
    ydot[0] = -k1*y[0] + k2*y[1]*y[2];
    ydot[2] = k3 * y[1]*y[1];
    ydot[1] = -ydot[0]-ydot[2];

    yout[0] = y[0]+y[1]+y[2];
}

/* The Jacobian matrix */
void jac(int *neq, double *t, double *y, int *ml, int *mu,
         double *pd, int *nrowpd, double *yout, int *ip)
{
    pd[0] = -k1;
    pd[1] = k1;
    pd[2] = 0.0;
    pd[(*nrowpd)] = k2*y[2];
    pd[(*nrowpd) + 1] = -k2*y[2] - 2*k3*y[1];
    pd[(*nrowpd) + 2] = 2*k3*y[1];
    pd[(*nrowpd)*2] = k2*y[1];
    pd[2*(nrowpd) + 1] = -k2 * y[1];
    pd[2*(nrowpd) + 2] = 0.0;
}

/* END file mymod.c */
```

## 2 Running ODE models implemented in compiled code

To run the C implementation of the model it must first be compiled. This can be done using the R `system` statement

```
system("R CMD SHLIB mymod.c")
```

which will create the file `mymod.dll` on Windows, or `mymod.so` on other platforms.

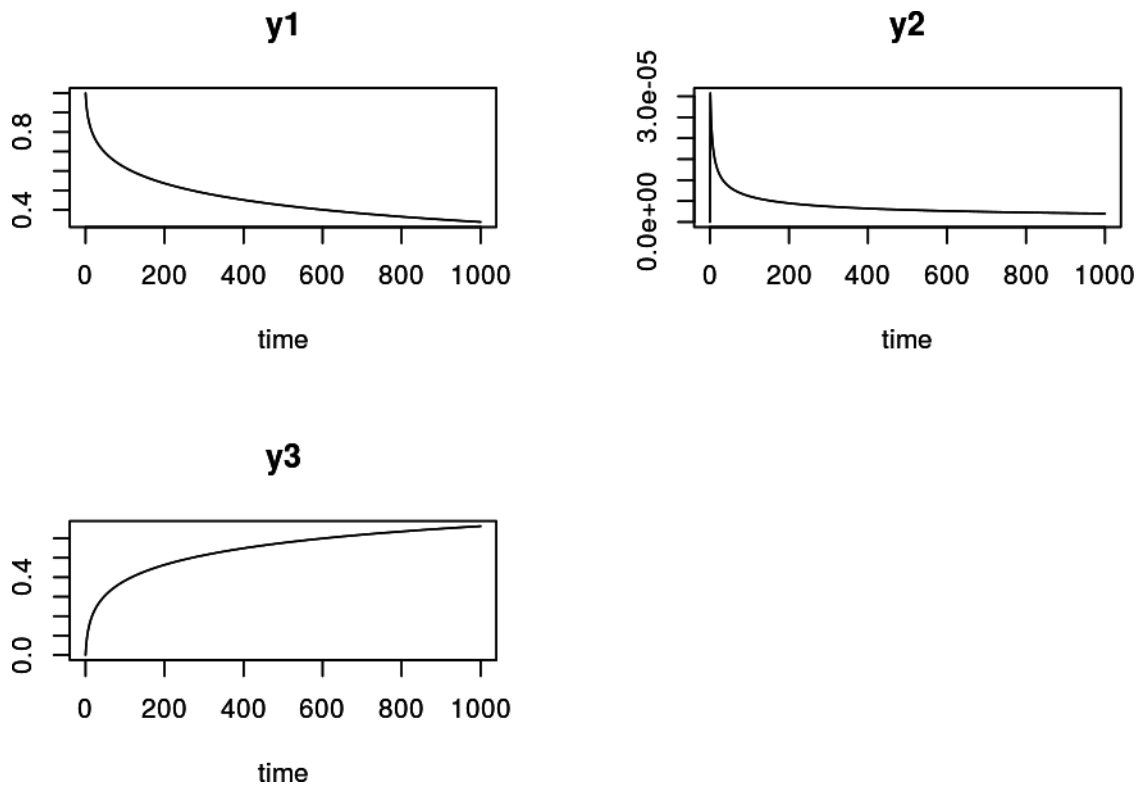
We can then load the compiled model with

```
dyn.load(paste("mymod", .Platform$dynlib.ext, sep = ""))
```

The model can now be run as follows:

```
parms <- c(k1 = 0.04, k2 = 1e4, k3=3e7)
Y      <- c(y1 = 1.0, y2 = 0.0, y3 = 0.0)
times  <- seq(0, 0.1*10^4, length.out = 1000)

out <- ode(Y, times, func = "derivs", parms = parms,
           jacfunc = "jac", dllname = "mymod",
           initfunc = "initmod", nout = 1, outnames = "Sum")
plot(out, which = 1:3)
```



We can determine the speed improvement as a function of desired output time points using the `microbenchmark` package

```

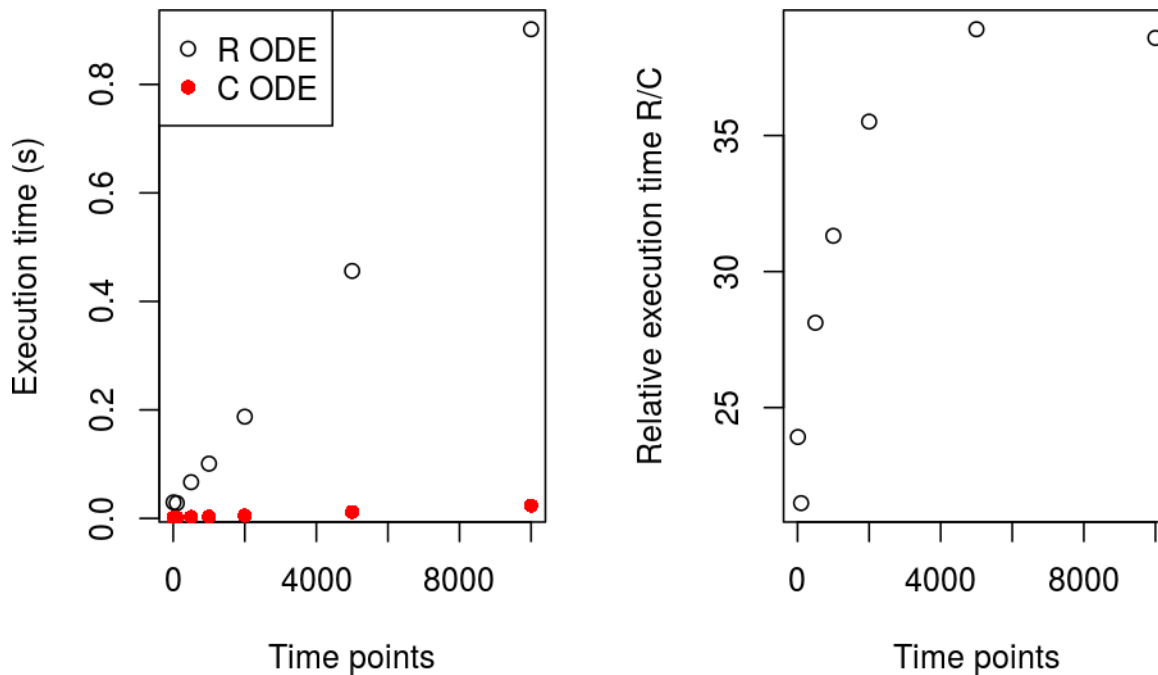
timepoints <- c(10,100, 500, 1000,2000, 5000, 10000)

solve_time <- lapply(timepoints, function(x){
  times <- seq(0, 0.1*10^4, length.out = x)
  comp <- microbenchmark::microbenchmark(
    C = ode(Y, times, func = "derivs", parms = parms,
            jacfunc = "jac", dllname = "mymod",
            initfunc = "initmod", nout = 1, outnames = "Sum"),
    R = ode(Y, times, modelR, parms = parms, jacfunc = jacR), times=10
  )
  return(comp)
})

#saveRDS(solve_time, "examples/Soetaert-r-vs-c-solve_time.RDS")

solve_speedup <- sapply(solve_time, function(x)mean(x$time[x$expr=="R"]/x$time[x$expr=="C"]))
solve_R <- sapply(solve_time, function(x)mean(x$time[x$expr=="R"])/1e9))
solve_C <- sapply(solve_time, function(x)mean(x$time[x$expr=="C"])/1e9))
#pdf("examples/c-vs-r-solve.pdf")
par(mfrow=c(1,2))
plot(timepoints, solve_R, type="p", ylab = "Execution time (s)", xlab = "Time points")
points(timepoints, solve_C, col="red", type="p", pch=c(16))
legend("topleft", legend = c("R ODE", "C ODE"), pch = c(1,16), col= c("black", "red"))
plot(rep(timepoints, each=1), solve_speedup, type="p", ylab = "Relative execution time R/C", xlab = "Time points")

```




```
#dev.off()
```

which shows us that the C version of the models is solved 21-39 times faster than the R version.

### 3 Conducting inference with compiled models

We can now simulate a dataset to demonstrate the inference procedure:

```
set.seed(143)
#force include the first time-point (t=0)
obs <- as.data.frame(outR[c(1:100,runif(100, 0, nrow(outR))),])
obs <- obs[order(obs$time),]
par(mfrow=c(1,3))
plot(obs$time, obs[,2], type='o')
plot(obs$time, obs[,3], type='o')
plot(obs$time, obs[,4], type='o')
```



### 4 Defining an observation model and parameters for inference

For simplicity we assume a normal log-likelihood for these data

$$\ell(\mathcal{Y}|\theta) = \sum_i \sum_t \ln \left( \frac{1}{\sigma_{obs} \sqrt{2\pi}} \exp \left( -\frac{(\tilde{y}_{t,i} - (y_{t,i}))^2}{2\sigma_{obs}^2} \right) \right) \quad (1)$$

where  $\tilde{y}_{t,i}$  are the observations, and  $y_{t,i}$  are the predictions of the DE model given the current MCMC sample of the parameters  $\theta$ .

```
# the observation model
obs_model <- function(data, sim.data, samp){

  llik.y1 <- sum(dnorm(obs[,2], mean = sim.data[,2], sd = samp[['sd.y1']], log = TRUE))
```

```

llik.y2 <- sum(dnorm(obs[,3], mean = sim.data[,3], sd = samp[['sd.y2']], log = TRUE))
llik.y3 <- sum(dnorm(obs[,4], mean = sim.data[,4], sd = samp[['sd.y3']], log = TRUE))
return(llik.y1 + llik.y2 + llik.y3)
}

```

We declare the DE model parameter  $r$ , assign a prior  $r \sim \mathcal{N}(0,1)$  and a random walk sampler with a Normal kernel (`samp.type="rw"`) and proposal variance of 0.005. Similarly, we declare  $K \sim \ln \mathcal{N}(1,1)$  and  $\ln(\sigma_{obs}^2) \sim \mathcal{N}(0,1)$ . Note that we are using the asymmetric uniform proposal distribution  $\mathcal{U}(\frac{a}{b}\theta^{(k)}, \frac{b}{a}\theta^{(k)})$  for the variance parameter (`samp.type="rw-unif"`), as this ensures strictly positive proposals.

```

library(deBInfer)
k1 <- debinfer_par(name = "k1", var.type = "de", fixed = FALSE,
  value = 0.05, prior = "lnorm", hypers = list(meanlog = 0, sdlog = 1),
  prop.var = c(998,1000), samp.type = "rw-unif")

k2 <- debinfer_par(name = "k2", var.type = "de", fixed = FALSE,
  value = 8000, prior = "norm", hypers = list(mean = 5000, sd = 1000),
  prop.var = 500, samp.type="rw")
k3 <- debinfer_par(name = "k3", var.type = "de", fixed = FALSE,
  value = 3e7, prior = "norm", hypers = list(mean = 1e7, sd = 5e6),
  prop.var = 5e7, samp.type="rw")

sd.y1 <- debinfer_par(name = "sd.y1", var.type = "obs", fixed = FALSE,
  value = 0.05, prior = "lnorm", hypers = list(meanlog = 0, sdlog = 1),
  prop.var = c(3,4), samp.type = "rw-unif")
sd.y2 <- debinfer_par(name = "sd.y2", var.type = "obs", fixed = FALSE,
  value = 0.05, prior = "lnorm", hypers = list(meanlog = 0, sdlog = 1),
  prop.var = c(3,4), samp.type = "rw-unif")
sd.y3 <- debinfer_par(name = "sd.y3", var.type = "obs", fixed = FALSE,
  value = 0.05, prior = "lnorm", hypers = list(meanlog = 0, sdlog = 1),
  prop.var = c(3,4), samp.type = "rw-unif")

#Lastly, we provide an initial value $N_0=0.1$ for the DE:

y1 <- debinfer_par(name = "y1", var.type = "init", fixed = TRUE, value = 1)
y2 <- debinfer_par(name = "y2", var.type = "init", fixed = TRUE, value = 0)
y3 <- debinfer_par(name = "y3", var.type = "init", fixed = TRUE, value = 0)

#The declared parameters are then collated using the `setup_debinfer` function. Note that for models wi

mcmc.pars <- setup_debinfer(k1, k2, k3, sd.y1, sd.y2, sd.y3, y1, y2, y3)

```

## 5 Conduct inference

Finally we use deBInfer to estimate the parameters of the original model

```

# do inference with deBInfer
# MCMC iterations
iter <- 1000
# inference call

```

```
Rt <- system.time(mcmc_samples <- de_mcmc(N = iter, data = obs, de.model = modelR,
  obs.model = obs_model, all.params = mcmc.pars,
  Tmax = max(obs$time), data.times = obs$time, cnt = 500,
  plot = FALSE, solver = "ode", verbose.mcmc = FALSE))
```

```
## Order of initial conditions is y1, y2, y3
```

```
## initial posterior probability = 568.058099497889
```

```
Ct <- system.time(mcmc_samplesC <- de_mcmc(N = iter, data = obs, de.model = "derivs",
  jacfunc = "jac", dllname = "mymod",
  initfunc = "initmod", nout = 1, outnames = "Sum",
  obs.model = obs_model, all.params = mcmc.pars,
  Tmax = max(obs$time), data.times = obs$time, cnt = 1250,
  plot = FALSE, solver = "ode", verbose.mcmc = FALSE))
```

```
## Order of initial conditions is y1, y2, y3
```

```
## initial posterior probability = 568.058099497889
```

```
print(Rt)
```

```
##      user  system elapsed
##  94.050   0.028  94.157
```

```
print(Ct)
```

```
##      user  system elapsed
##  13.429   0.004  13.444
```

```
#coda::rejectionRate(mcmc_samplesC$samples)
#plot(mcmc_samplesC)
```

, and simulate posterior trajectories.

```
system.time(post_traj <- post_sim(mcmc_samples, n=100, times=0:8000, burnin=100, output = 'all', prob =
```

```
##      user  system elapsed
##  80.080   0.104  80.226
```

```
system.time(post_trajC <- post_sim(mcmc_samplesC, n=100, times=0:8000, burnin=100, output = 'all', prob =
```


```
##      user  system elapsed
##  11.697   0.108  11.812
```



```

par(mfrow=c(1,3))
plot(post_traj$time, post_traj$median$y1, type='l', lwd=2)
lines(post_trajC$time, post_trajC$median$y1, col="red", lty=3, lwd=4)
points(obs$time, obs[,2])
plot(post_traj$time, post_traj$median$y2, type='l')
lines(post_trajC$time, post_trajC$median$y2, col="red", lty=3, lwd=4)
points(obs$time, obs[,3])
plot(post_traj$time, post_traj$median$y3, type='l')
lines(post_trajC$time, post_trajC$median$y3, col="red", lty=3, lwd=4)
points(obs$time, obs[,4])

```



We can that inference and posterior simulations using the compiled model are substantially faster. And this speed up is consistent across a range of MCMC iterations:

```

iter_reps <- rep(c(100,1000,2000,5000,10000),3)
timings <- matrix(nrow=length(iter_reps),ncol=2)

for (i in seq_along(iter_reps)){

  timings[i,1] <- system.time(mcmc_samples <- de_mcmc(N = iter_reps[i], data = obs, de.model = modelR,
    obs.model = obs_model, all.params = mcmc.pars,
    Tmax = max(obs$time), data.times = obs$time, cnt = 500,
    plot = FALSE, solver = "ode", verbose.mcmc = FALSE))[3]
  timings[i,2] <- system.time(mcmc_samplesC <- de_mcmc(N = iter_reps[i], data = obs, de.model = "derivs",
    jacfunc = "jac", dllname = "mymod",
    initfunc = "initmod", nout = 1, outnames = "Sum",
    obs.model = obs_model, all.params = mcmc.pars,
    Tmax = max(obs$time), data.times = obs$time, cnt = 1250,
    plot = FALSE, solver = "ode", verbose.mcmc = FALSE))[3]
}

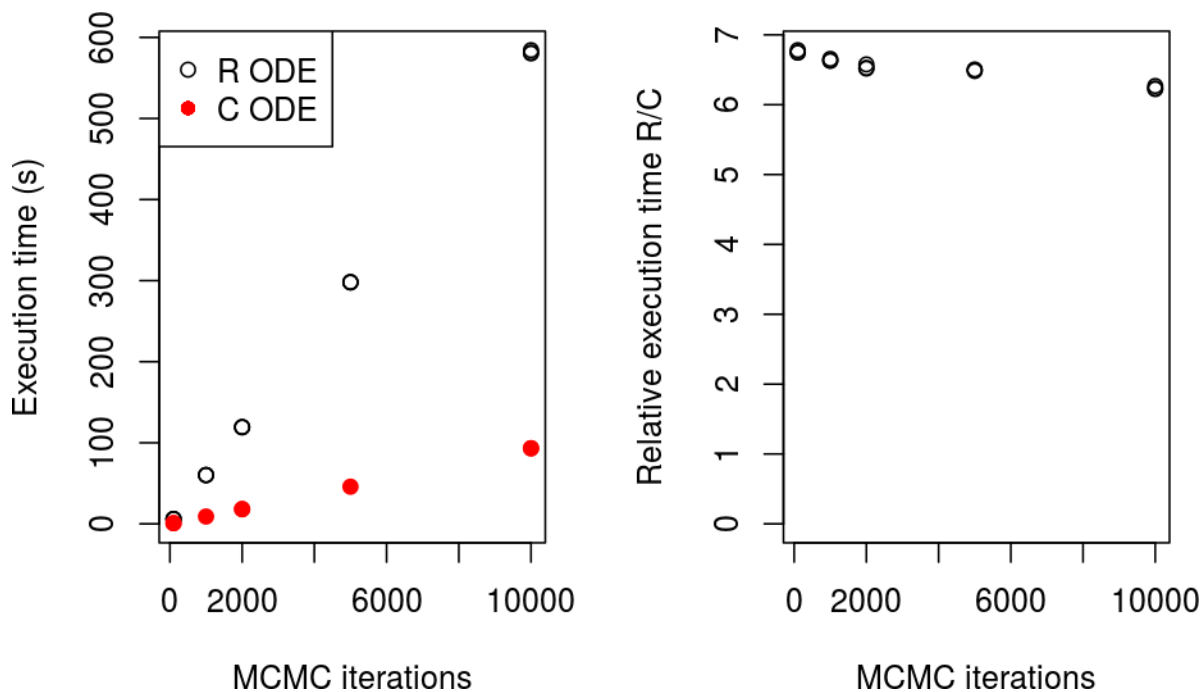
```

```

}
#saveRDS(timings, "examples/Soetaert-r-vs-c-timings.RDS")

iter_reps <- rep(c(100,1000,2000,5000,10000),3)
if(!exists("timings")) timings <- readRDS("Soetaert-r-vs-c-timings.RDS")
#pdf("examples/c-vs-r-iters.pdf")
par(mfrow=c(1,2))
plot(iter_reps, timings[,1], ylim=c(0, max(timings[,1])), type="p", ylab = "Execution time (s)", xlab = "MCMC iterations")
points(iter_reps, timings[,2], col="red", type="p", pch=c(1,16))
legend("topleft", legend = c("R ODE", "C ODE"), pch = c(1,16), col= c("black", "red"))
plot(iter_reps, timings[,1]/timings[,2], ylim=c(0, max(timings[,1]/timings[,2])), type="p", ylab = "Relative execution time R/C", xlab = "MCMC iterations")

```



```
#dev.off()
```

Boersch-Supan, P. & Johnson, L. (2016). deBinfer: Bayesian inference for dynamical models of biological systems. *arXiv*, **1605.00021**.

Soetaert, K., Petzoldt, T. & Setzer, R.W. (2009). *R package desolve, writing code in compiled languages*.

Soetaert, K., Petzoldt, T. & Setzer, R.W. (2010). Solving differential equations in R: Package deSolve. *Journal of Statistical Software*, **33**, 1–25.